

# ADVANCED PARALLEL COMPUTING 2017

## LECTURE 05 - SYNCHRONIZATION 2

Holger Fröning  
[holger.froening@ziti.uni-heidelberg.de](mailto:holger.froening@ziti.uni-heidelberg.de)  
Institute of Computer Engineering  
Ruprecht-Karls University of Heidelberg

*Some material by Falsafi, Hardavellas, Nowatzky of EPFL, Northwestern, CMU*

# BARRIERS



# WHY DO WE NEED BARRIERS?

What is a barrier?

Any participating thread has to wait until all threads have reached the barrier

Why do we need such coarse-grain (collective) synchronization?

Any kind of lock-step computation

Physics simulation computation


Divide up the computation of each time step into N independent pieces

Each time step: compute independently, synchronize

Essential for BSP-like programming

Stencil codes

Six-step FFT method



Barrier-based Merge-Sort

# BARRIERS VS. LOCKS

## Lock

Threads are waiting because they have something to do

A lock master is usually not viable

- No thread/core may be free to take over this part

- Contention would quickly turn the master into a bottleneck

## Barrier

Threads are rather waiting because they have finished their work and wait for their external dependencies to be resolved

A barrier master can be viable

- At least one thread is waiting anyway

# GLOBAL SYNCHRONIZATION BARRIER

At a barrier, all threads wait until all other threads have reached it

# Strawman implementation (wrong)

# So what's wrong?

# Race condition

After resetting counter to P, before all threads are notified, another thread could again decrement counter


```
(shared) int count = P;

void strawman_barrier ()
{
    if ( fetch_and_dec ( &count ) == 1 )
        count = P;
    else
        while ( count != P );
        //spin wait
}
```

```

...
do_work_1 ();
strawman_barrier ();
do_work_2 ();
strawman_barrier ();
...

```



# FUNCTIONALLY CORRECT BARRIER

## “SENSE-REVERSING”

```
(shared) int count = P;
(shared) bool sense = true;
(local) bool local_sense = true;

void sense_barrier ()
{
    // each processor toggles its own sense
    local_sense = !local_sense;

    if ( fetch_and_dec ( &count ) == 1 )
        count = P;
        // last processor toggles global sense
        sense = local_sense;
    else
        while ( sense != local_sense ); //spin wait
}
```

# ISSUES WITH CENTRALIZED BARRIERS

Single counter makes the sense-reversing barrier a centralized barrier

Highly suitable for cache-coherent shared memory systems

Cache coherence & polling on sense:  
uniform notification time (hopefully)


Problem with centralized barrier

All processors must increment the counter

Each RMW is a serialized coherence action, i.e. a cache miss

$O(n)$  threads arrive simultaneously, slow for lots of processors

=> Memory contention





# COMBINING TREE BARRIER

Combining tree barrier: build a  $\log_k(n)$  height tree of counters (one per cache block)

$k$  = radix

Each thread coordinates with  $(k-1)$  other threads (i.e., by thread ID)

Last of the  $k$  threads coordinates with next higher node in tree (no tournament - statically determined)

As many coordination addresses are used, misses are not serialized ( $O(\log n)$  in best case)


Spread memory accesses across multiple barriers

=> Reduce contention

What is faster, decrement a single location or visiting a logarithmic number of barriers?

=> At the cost of increased latency

Eventually





# OTHER BARRIER IMPLEMENTATIONS - SW

Separate barrier operation into two phases: arrival and release

Various optimizations



- Centralized or (software) combining tree

  - Different fan-outs, even different for arrival/release tree

- Dissemination for local spinning

- Tournament to avoid RMW

And these are only the software implementations



# OTHER BARRIER IMPLEMENTATIONS - HW

In principle: AND trees (possibly even hard-wired)

## Issues

- No one wants a second dedicated network

- Virtualization: support for more than one concurrent barrier

## Proposals & Implementations

- Syncbits associated with each cache block

  - J. Goodman et al., “Efficient Synchronization Primitives for Large- Scale Cache-Coherent Shared-Memory Multiprocessors”, ASPLOS1989

- E-Registers & synchronization units (32, memory-mapped) (Cray T3E)

  - S.L. Scott, “Synchronization and Communication in the T3E Multiprocessor”, ASPLOS1996

- Starving cache line requests (postpone responses until barrier condition met)

  - J. Sampson et al., “Exploiting Fine-Grained Data Parallelism with Chip Multiprocessors and Fast Barriers”, MICRO2006

- GBarrier for many-core CMPs (G-Line based dedicated network)

  - J. L. Abellán, J. Fernández and M. E. Acacio, "Efficient Hardware Barrier Synchronization in Many-Core CMPs", TPDS2012

# BARRIER PERFORMANCE

Barrier performance on Symmetry (shared-memory bus-based 80386 system, cache-based)

## Network transactions

Dissemination barrier:  $O(P \log P)$ , tournament & tree:  $O(P)$ , centralized & arrival tree:  $O(P)$  if broadcast support, potentially unbounded otherwise

## Critical path length

Assuming networks supports parallel transactions

Centralized:  $O(P)$ , others:  $O(\log P)$


## Space

Centralized: 1, tree:  $O(P)$ , others:  $O(P \log P)$

## Arrival tree

Tree barrier with a sense-reversing flag instead of a release tree

Lower slope than counter as  $N-1$  writes are cheaper than  $N$  RMWs



Mellor-Crummey, Scott, TOCS 1991

# BARRIER PERFORMANCE

Barrier over distributed shared memory  
(MEMSCALE)

No global coherence but local coherence domains

Support for remote loads and stores

No RMW


Tree-based: scalability, locality

Tournament: arrays of flags instead of RMWs

Dissemination: push instead of pull

=> 1k threads in 15usec

64P @IB: min. 20us



# OTHER ASPECTS

## Relaxations

Do you really need to synchronize?

Stencil codes with halo exchange: Probably at the cost of an increase in iterations

Training of deep learning: asynchronous weight exchange

## Overlap

Hide barrier latency with communication

Which model does that?

## Jitter

Large-scale SPMD programs might suffer from varying state  
=> increased synchronization time

# SUMMARY

## Spin Lock

Avoid „Coherence Storms“

Make use of atomic operations if available

Preferable Compare-and-Swap, or Fetch-and-Store

Remember the Bakery algorithm as SW alternative

List-based queuing locks usually outperform any other

In particular true for scalability

HW like QOLB: ROI too low

## Barriers

Avoid „Coherence Storms“

Avoiding atomics can be beneficial

Use combining tree to minimize the critical path length, and dissemination to leverage locality effects

Treat arrival and release structures differently

HW structures

ROI too low for shared memory

Potentially reasonable for message passing



# PROJECT WORK PROPOSALS





# CONSTRUCTION AREA

## NVIDIA Fermi-class GPU

2 per server, 8 servers

## NVIDIA Kepler-class GPU

16 per server (1x)

1 per server (8x)

Embedded platform (Tegra TK1, TX1)

## NVIDIA Pascal-class GPU

2 per server (1x)

## ARM Mali GPU

Embedded platform

## ARM CPU

Tegra TK1, TX1

A couple more...

## Micron PicoComputing System

OpenCL-capable FPGA+HMC


## Intel X86

Sandy Bridge (12 core, 8x)

Ivy Bridge (16 core, 1x)

## Intel MIC

Possibly, but have to check...



# ANALYZING DYNAMIC NUMA EFFECTS: LOAD VERSUS LATENCY

Static NUMA effects: memory access latency variations for an unloaded system

Dynamic NUMA effects: memory access latency variations for a loaded vs unloaded system

GEMM as a simple load

Architectures

Multi-GPU

x86

ARM

Goal: understand where this is an issue

# TESTING GPU MEMORY MODELS

GPU memory models are weak and only informally defined

Approach

- Develop a set of tests (litmus tests)

- Evaluate different GPU architectures

Complex task, but plenty of related work

Goal: test memory models for fun (and profit)

# IMPLEMENTING A SCALABLE DATA ANALYSIS FOR EXASCALE APPLICATIONS

Problem: we have more data than we can analyze in reasonable time

- Python and bash scripts

- Traces for MPI Exascale proxy applications: ~1/2TB analyzed, >25TB left

Open questions

- Do we need ordering?

- What is the impact of matching on overall performance?

Spark would be an interesting tool

Impact: huge

- We're already best paper finalist for ISC2017 paper (see website)

# EXPLORING UNIFIED MEMORY: PERFORMANCE, ENERGY AND (MAYBE) SCHEDULING IMPLICATIONS

“Downgrade” some CUDA SDK apps to Unified Memory

Replace `cudaMalloc` + `cudaMemcpy` with `cudaMallocManaged`

Explore performance and energy on NVIDIA Pascal GPUs

Option: Interact with some multi-GPU guys

Design micro-benchmarks that highlight performance and fallacies

Optionally: prefetch instructions

Goal: understanding + micro-benchmarks



# GLOBALLY VALID TIME MEASUREMENT IN CUDA KERNELS

Problem: CUDA GPUs have only a local time (one counter per SM)

To understand scheduling effects better, a global time is required

Task: Design a method that synchronizes the different clock counters so that the notion of a global time can be established

Architecture

NVIDIA GPU

# TRACING SCHEDULING DECISIONS FOR CUDA APPLICATIONS

Problem: CUDA scheduling is mostly unknown

Known: greedy-then-oldest, no preemption

Goal: more understanding

Approach: track CTA start/end time and physical SM ID

Explore how scheduling can be influenced (streams, kernel sharding, kernel fission)

Explore co-scheduling of multiple kernels

Dependencies: previous



# IMPACT OF MATRIX MULTIPLY OPTIMIZATIONS ON PORTABILITY

Conventional wisdom: optimizations hurt portability

## Approach

- Implement different optimizations for an Open-CL matrix multiply

- Explore their impact on different processors

## Architectures

- NVIDIA {Fermi, Kepler, Pascal}, embedded-class & server-class

- x86 CPU

- ARM CPU

- ARM GPU (Mali)

Possibly revert to CUDA

# DISSECTING GPU MEMORY HIERARCHY

Recent work has shown that GPU's have huge TLB issues

<http://dl.acm.org/citation.cfm?doid=3076113.3076115>

In general, understanding GPUs in more detail is important

<https://arxiv.org/pdf/1509.02308.pdf>

[http://www.comp.hkbu.edu.hk/~chxw/gpu\\_benchmark.html](http://www.comp.hkbu.edu.hk/~chxw/gpu_benchmark.html)

Goal: Adapt existing benchmarks and possibly come up with new ones

Focus on TLB & address translation

# OPTIMIZE LIST-BASED SEARCH FOR MPI MATCHING

MPI Matching is an important problem in HPC communication stacks

- Two lists: unexpected messages, and posted receives

- Ordering guarantees require sequential searching

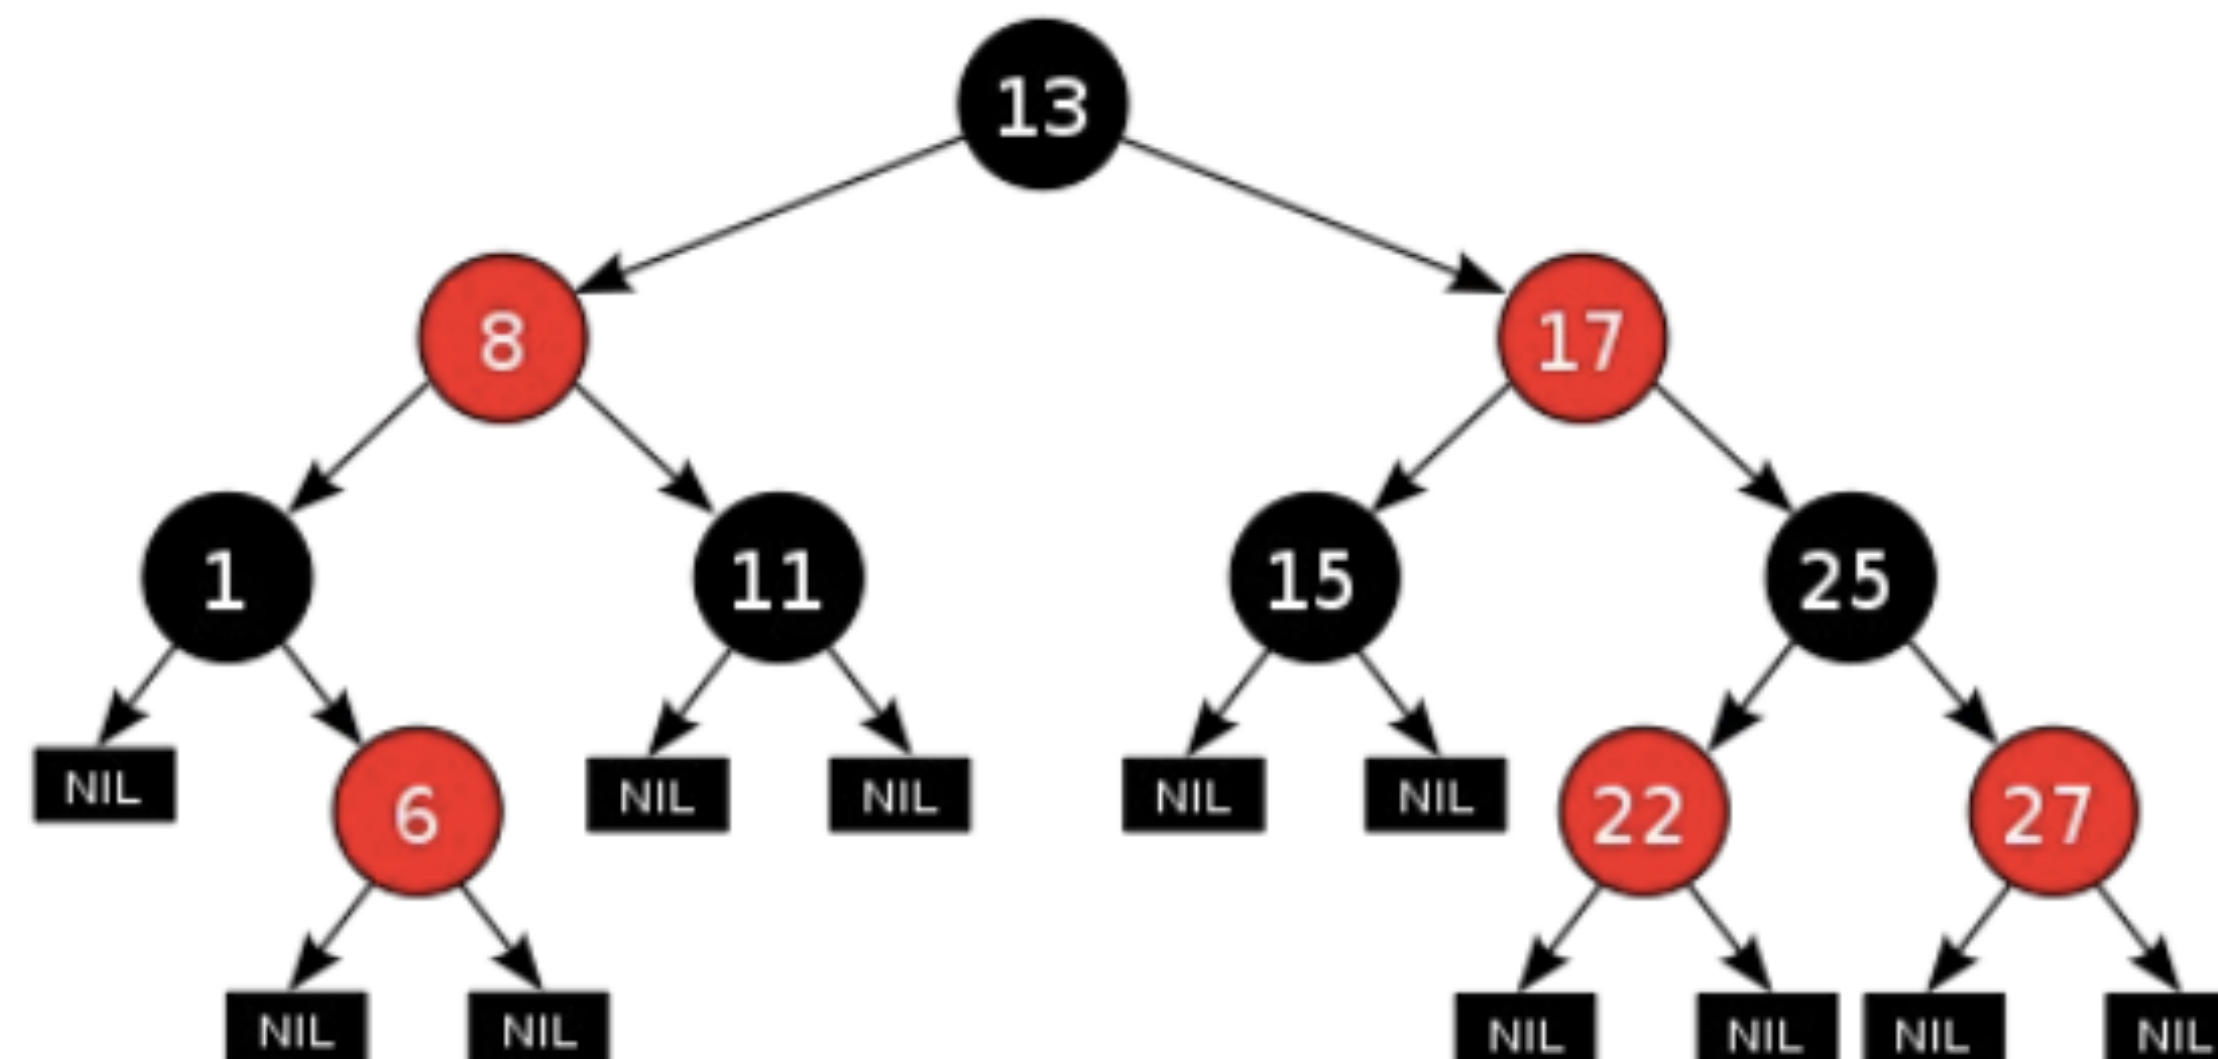
MPI implementations achieve about 25M matches per second, but vary significantly with list length

- Can we understand the different optimizations?

- Can we do better?

# OPTIMIZING A RED-BLACK TREE USING LOCKING

Red-Black Tree (RBT) is a form of self-balancing binary search tree  
Complex data structure to explore the impact of different locking schemes and methods



# POCL

Portable Computing Language (pocl) aims to become a MIT-licensed open source implementation of the OpenCL standard which can be easily adapted for new targets and devices, both for homogeneous CPU and heterogeneous GPUs/accelerators.

<https://github.com/pocl/pocl>

## Goal

- Explore portability

- Possibly using different optimizations of a matrix multiply

## Architectures

- x86 CPU

- ARM CPU?

- NVIDIA GPU

# MULTI-GPU CONVOLUTION BASED ON THE CUDA SDK VERSION

Addressing the permanent lack of multi-GPU benchmarks

Convolution is an important operation for ML

## Approach

Start with the CUDA SDK single-GPU version

Implement and optimize a multi-GPU version

Optionally: replace <convolution> with any other app

# NOTES

Anyone eager to work with architecture simulators?

Make up your mind - we're open for proposals!

Think about aligning this project assignment with future work

Seminar talk, project work, MSc thesis