

ADVANCED PARALLEL COMPUTING 2017

LECTURE 06 - TRANSACTIONAL MEMORY

Holger Fröning
holger.froening@ziti.uni-heidelberg.de
Institute of Computer Engineering
Ruprecht-Karls University of Heidelberg

Some material by Falsafi, Hardavellas, Nowatzky of EPFL, Northwestern, CMU

**LOCKING: PRICE OF GUARANTEED
CORRECTNESS**

LOCKS

Getting locking right

- Taking too few locks

- Taking too many locks

- Taking the wrong locks

- Taking the locks in wrong order

- Freeing locks on error

Locking is expensive

- Coherence actions

- Wait/Transfer/Release

Locking is pessimistic



AN EXAMPLE

A concurrent data structure called “set”

Operations on this data structure

Concurrency

```
insert (set s, key k)
lookup (set s, key k)
delete (set s, key k)
```

Insert/Delete/Lookup should be mutually exclusive

Now, support: `transfer (set s1, set s2, key k)`

Requirement: k has always to be stored in one of both sets

Never both or neither

Fine grain locking: lock elements, deadlocks!

Coarse grain locking: lock data structure, less concurrency!

AN EXAMPLE

Even with coarse grain locking

Breaks abstraction: exposes internal lock

Deadlock: which set's lock to grab first?

An abstract but ideal solution

```
void transfer (set s1, set s2, key k)
{
    atomic {
        delete (s1, k)
        insert (s2, k)
    }
}
```

Where “atomic” has

Simplicity of coarse grain locking

Concurrency of fine grain locking

Without fine-grain locking overheads

=> The promise of Transactional Memory

TRANSACTIONAL MEMORY

Region that executes serially (isolated/atomic)

Inspired by database transactions, but different

Implementation: **speculative execution**

Serialize only on dynamic conflicts (eager or lazy)

E.g., when key manipulated by different threads

Partly overcomes the granularity/complexity tradeoff

Avoid conservative serialization of locking

Mutual exclusion: **Optimistic** instead of pessimistic

HOT TOPIC / GARTNER'S HYPE CYCLE

Pioneering work

First idea by Lomet, 1977

Herlihy+, ISCA1993

Speculative locking

E.g. Rajwar+, MICRO2001 & ASPLOS 2002

Software Transactional Memory

E.g. Herlihy+, PODC2003; Harris+, OOPSLA 2003; ...

Hardware Transactional Memory

E.g. Hammond+, ISCA2004, ASPLOS2004; UTM (HPCA2005);
VTM (ISCA2006); LogTM (HPCA2006); ...


HW/SW Hybrids...

Commercial implementations

IBM BlueGene/Q (2011), Intel Haswell (2012)

Lots of TM papers in the recent years

300+ citations in „Transactional Memory“, 2nd Ed., 2010



Source: wikipedia.com

SPECULATIVE LOCKING

Correctly synchronizing a program with locks is hard

Fine-grain locking

Difficult to program

High overhead

Coarse-grain locking

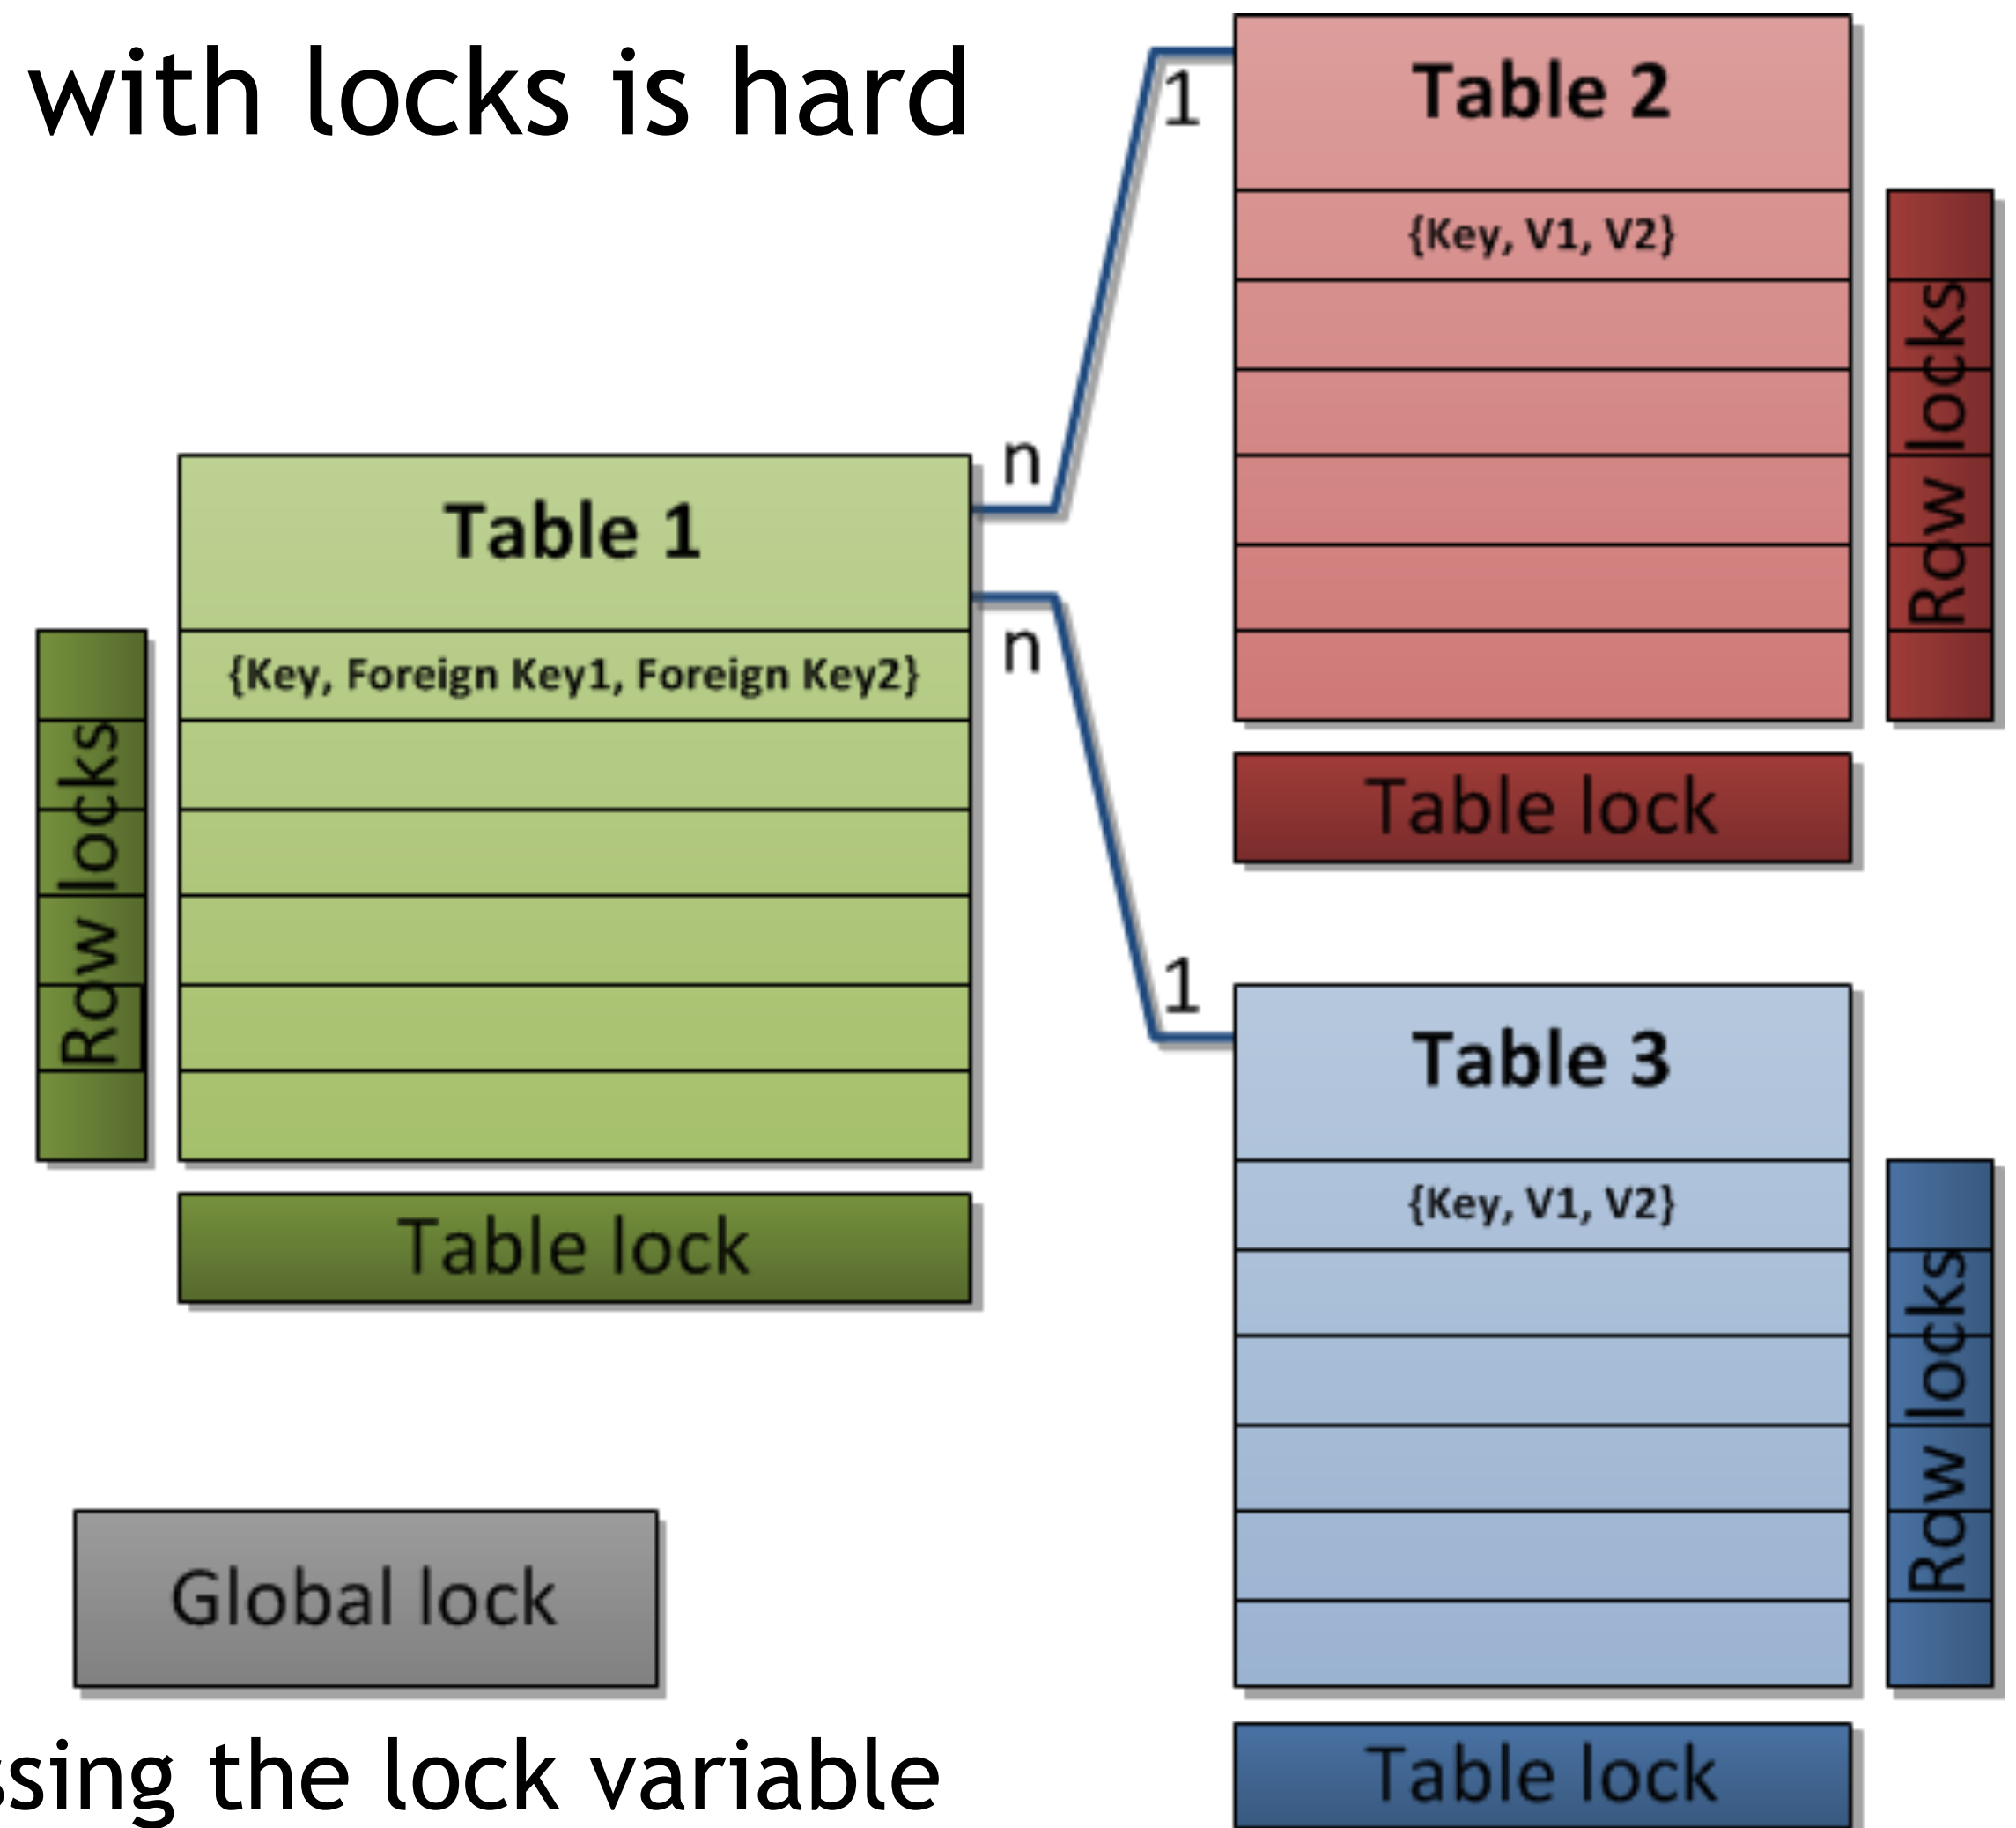
Poor performance

Even worse scalability

But, concurrent critical sections usually access disjoint data

So, they could run in parallel...

... except that they conflict on accessing the lock variable



SPECULATIVE LOCK ELISION

Speculatively execute critical sections in parallel

Key idea: Detect and elide lock access

- Upon a lock acquire, don't actually acquire the lock

- Checkpoint the processor state

- Run critical sections in parallel

- Detect conflicting data accesses via coherence protocol

- Any invalidates before the lock release cause rollback

- Then retry by acquiring the lock normally

Advantages

- No locking overhead, since lock is not actually acquired

- Allows concurrent execution of non-conflicting critical sections

How to detect critical sections?

- Detect temporally silent stores

INTRODUCTION TO TRANSACTIONAL MEMORY

TM'S BIG IDEAS

Big Idea #1: no locks, just shared data

Big Idea #2: optimistic (speculative) concurrency

Execute critical sections speculatively, abort on conflict

“Better to beg for forgiveness than to ask for permission”

```
struct account_t { int balance; };  
shared struct account_t a[MAX];  
int fromID, toID, amount;  
  
begin_transaction();  
if (a[id_from].balance >= amount) {  
    a[fromID].balance -= amount;  
    a[toID].balance += amount;  
}  
end_transaction();
```

TM'S READ/WRITE SETS

Read set: set of shared addresses critical sections reads

Example: `a[37].balance, a[241].balance`

Write set: set of shared addresses critical section writes

Example: `a[37].balance, a[241].balance`

```
struct account_t { int balance; };  
shared struct account_t a[MAX];  
int fromID, toID, amount;  
  
begin_transaction();  
if (a[id_from].balance >= amount) {  
    a[fromID].balance -= amount;  
    a[toID].balance += amount;  
}  
end_transaction();
```

TM'S BEGIN

`begin_transaction()`

Take a local register snapshot

Begin locally tracking read set (remember addresses you've read), see if anyone else is trying to write it

Locally buffer all your writes (invisible to other processes)

=> Local actions only: no lock acquire

```
struct account_t { int balance; };
shared struct account_t a[MAX];
int fromID, toID, amount;

begin_transaction();
if (a[id_from].balance >= amount) {
    a[fromID].balance -= amount;
    a[toID].balance += amount;
}
end_transaction();
```


TM'S END

`end_transaction()`

Check read set: is all your data still valid (no writes to any)?

Yes: commit transaction by committing writes

No: abort transaction by restoring checkpoint

Writes have to appear atomically with regard to other threads!

```
struct account_t { int balance; };  
shared struct account_t a[MAX];  
int fromID, toID, amount;  
  
begin_transaction();  
if (a[id_from].balance >= amount) {  
    a[fromID].balance -= amount;  
    a[toID].balance += amount;  
}  
end_transaction();
```

TRANSACTIONAL EXECUTION (CONFLICT)

Thread 0

```
fromID = 241;
toID = 37;

begin_transaction();
if(a[241].balance > 100) {
    ...
    ...
    // detect write to
    //   a[241].balance
    // abort, rollback and retry
```

Thread 1

```
fromID = 37;
toID = 241;

begin_transaction();
if(a[37].balance > 100) {
    a[37].balance -= amount;
    a[241].balance += amount;
}
end_transaction();
// no writes to a[241].balance
// no writes to a[37].balance
// commit and proceed
```

TRANSACTIONAL EXECUTION (MORE LIKELY)

Thread 0

```
fromID = 241;
toID = 37;

begin_transaction();
if(a[241].balance > 100) {
    a[241].balance -= amount;
    a[37].balance += amount;
}
end_transaction();
// no write to a[241].balance
// no write to a[37].balance
// commit and proceed
```

Thread 1

```
fromID = 450;
toID = 118;

begin_transaction();
if(a[450].balance > 100) {
    a[450].balance -= amount;
    a[118].balance += amount;
}
end_transaction();
// no write to a[450].balance
// no write to a[118].balance
// commit and proceed
```

Critical sections execute completely in parallel, as there is no conflict detected.

IMPLEMENTATION DESIGN SPACE

Four main components

1. Version management

Logging/buffering

Register and memory accesses

2. Conflict detection

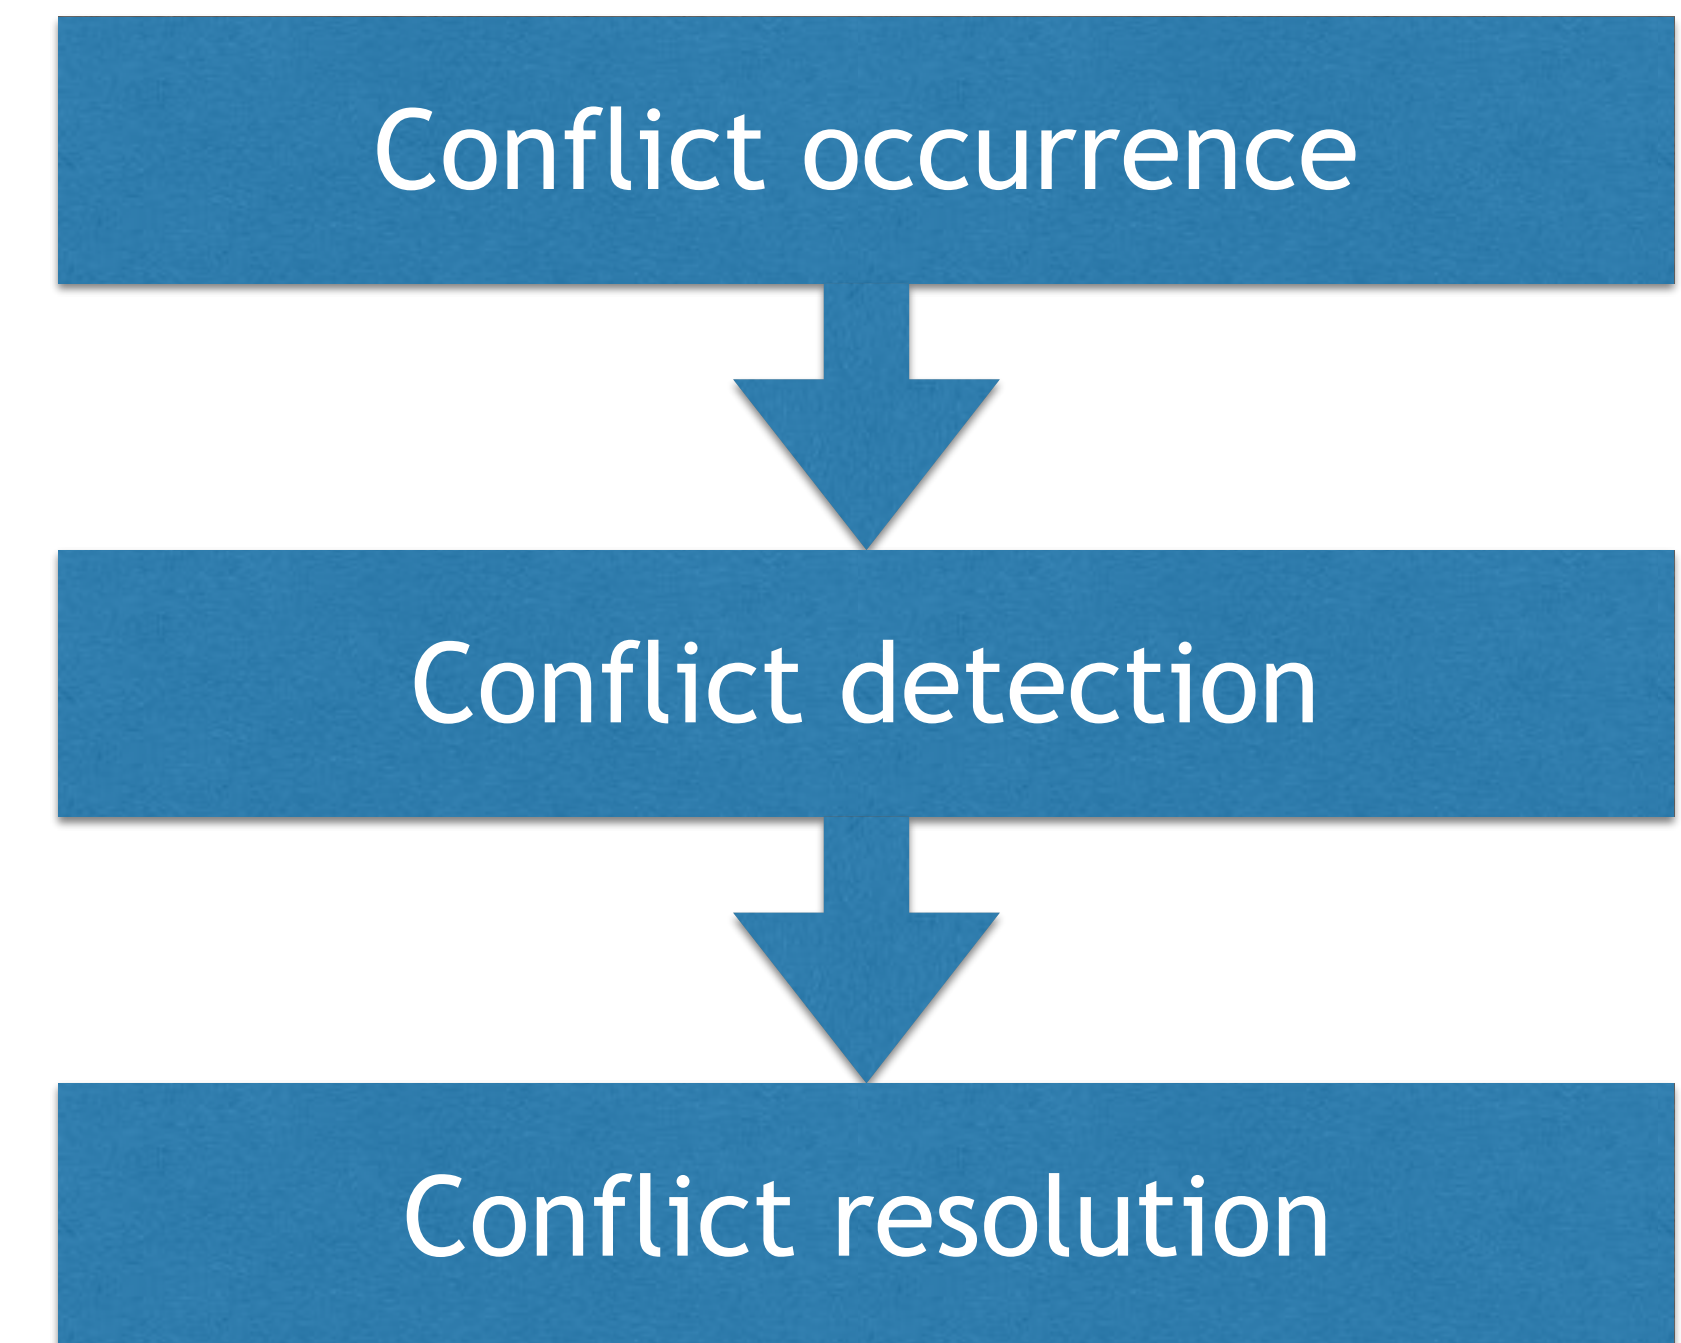
Two accesses to a single location,
at least one is a write

3. Abort/Rollback

4. Commit

Various implementation approaches

Hardware TM, Software TM, Accelerated STM, Hybrids



VERSION MANAGEMENT: PRESERVING REGISTER VALUES

Begin transaction

- Take register checkpoint

Commit transaction

- Free register checkpoint

Abort transaction

- Restore register checkpoint

Register checkpoint includes all registers

- Including control registers like PC, stack pointers, etc.

VERSION MANAGEMENT FOR MEMORY

Lazy version management

Store

Put all writes into write table

Load

If address in write table, read value from write table

“Read own write early”

Otherwise, read from memory

Commit transaction

Write all entries from write table to memory, clear it

Abort transaction

Clear write table (fast)

Eager version management

Store

If address not in write set:

1. Read old value and put it into write log
2. Add address to write set

Write stores directly to memory

Load

Read directly from memory (fast)

Commit transaction

Nothing (fast)

Abort transaction

Traverse log, write logged values back into memory (slow)

CONFLICT DETECTION

Lazy conflict detection

Store

Add address to write set (if not already present)

Load

Add address to read set (if not already present)

Commit transaction

For each address A in write set, and for each other thread T: if A is in T's read set, abort T's transaction

Eager conflict detection

Store

Add address A to write set (if not already present)

For each other thread T: if A is in T's write or read set, trigger conflict

Load

Add address to read set (if not already present)

For each other thread T: if A is in T's write set, trigger conflict

Conflict: abort either transaction

Commit transaction

If not already aborted, commit and clear sets

SOFTWARE TRANSACTIONAL MEMORY (STM)

Add extra software to perform
TM operations

Version management

- Software data structure for write
table (log)

- Eager or lazy

Conflict detection

- Software data structure (lock
table), mostly lazy

- Too high overhead for eager
implementations

- Granularity: object or block

Commit

- Need to ensure atomic update of all
states

- Grabs lots of locks, or a global
commit lock

Many possible implementations
and semantics

- Various (disjoint) goals like low
sequential overhead, or good
scalability, or strong progress
guarantees, ...

HARDWARE TRANSACTIONAL MEMORY (HTM)

Key idea: leverage invalidation-based coherence

Each cache block has read-only or read-write state

Coherence invariant:

Many read-only (shared) blocks, or
Single read-write block

Add two bits per cache block: read & write

Set on load/store during transactional execution

If another processor steals block from cache, abort

1. Read or write request to cache block with write bit set
2. Write request to cache block with read bit set

Low-overhead conflict detection

But only if all data fits into cache

Eviction: losing information associated with transaction

Implementations and semantics

ReadTx/WriteTx or atomic blocks (all memory accesses are implicitly transactional)

Automatic re-execution after aborts

Use of coherence or development of new memory systems for TM

Use of write-buffers?

HTM VS STM

HTM

Requires hardware (~~announced~~
~~available cancelled~~ available)

Simple for bounded case

Unbounded TM in hardware really complicated

Size: tracking conflicts after cache overflow

Duration: context switching during transactions

Cache block granularity for conflicts

STM

Here today (prototype compilers from Intel and others)

Generally weaker semantics

Slow (2x or more single-thread overhead)

Adds lots of instructions to memory operations

HYBRID TRANSACTIONAL MEMORY

Hardware-accelerated STM

Add special hardware tracking features

Under control of software

Can reduce STM overhead, but perhaps not enough

Hybrid HTM/STM

Use HTM mode most of the time

Revert to STM only on overflows and such

Getting the interaction right is actually really tricky

SOME CAVEATS

Although transactional memory really looks promising...

What if: read or write set bigger than cache?

What if: transaction gets scheduled out in the middle?

What if: transaction wants to do IO or syscalls?

Concurrency issues (see databases)

- Between transactions & between TM ops

- Linearizability (atomicity) & serializability (basic correctness condition)

How do we translate lock-based programs into TM programs?

- Replacing `lock_acquire` with `begin_transaction` doesn't always work

Several kinds of transaction semantics

- Are transactions atomic relative to code outside of transactions?

TRANSACTIONS ARE NOT CRITICAL SECTIONS

What's wrong with this code?

Thread 0

```
begin_transaction();  
flagA = true;  
while (!flagB);  
//update m  
end_transaction();
```

Thread 1

```
begin_transaction();  
while (!flagA);  
flagB = true;  
//update n  
end_transaction();
```

A less contrived example:

Thread 0

```
begin_transaction();  
...  
queueA->enqueue(val1);  
while (queueB->empty());  
//access queueB  
...  
end_transaction();
```

Thread 1

```
begin_transaction();  
...  
queueB->enqueue(val2);  
while (queueA->empty());  
//access queueA  
...  
end_transaction();
```

These are not necessarily programming errors...

MORE ISSUES: CONSISTENCY DURING TRANSACTIONS

Consistency during transactions not covered by DBs!

Strict serializability provides an intuitive model for the execution of committed transactions - nothing about failures

Nothing about the interactions of TM-/non-TM accesses

Inconsistent Reads

Example: eager VM and lazy CD => zombie or “doomed” transactions

User is responsible to do an “incremental validation”

Validating n locations requires n memory accesses; if for each access (total m): $O(n*m)$

Thread 0

```
do {
  StartTx();
  int tmp_1 = ReadTx(&x);

  int tmp_2 = ReadTx(&y);
  while (tmp_1 != tmp_2) { }
} while (!CommitTx());
```

Thread 1

```
do {
  StartTx();
  WriteTx(&x, 10);
  WriteTx(&y, 10);
} while (!CommitTx());
```

LCD: For each address A in write set, and for each other thread T : if A is in T 's read set, abort T 's transaction

IMPLEMENTATIONS

STM

Mainly two types

Fully validates: validate complete read set after each transactional read: more concurrency at a higher price

Global version number: avoid full validation by comparing version numbers at the end of the transaction (one of the best trade-offs)

Overhead: 1.5-120x single-thread execution time

Mostly dead respectively an extension to HTM => hybrid HTM

HTM

For two decades studied in academia

Now first commercial solutions available

Intel Transactional Synchronization Extensions (TSX): XACQUIRE/XRELEASE & XBEGIN/XEND

Intel's Haswell Broadwell, IBM's Blue Gene/Q, Oracle: future SPARC will include memory versioning, AMD: either goes its own way with ASF, or joins Intel's HLE/RTM

SUMMARY

Transactional Memory is a promising technique to address:

- Programming (synchronization) complexity

- Overhead associated with traditional locking

Caveats

- Transactional Memory != transactions in databases

- Transactions != critical sections

- Mixture of transactional and non-transactional traffic

- Pessimistic locking => probability of deadlocks

- Optimistic TM => probability of livelocks

SUMMARY: TM OVERVIEW

